

LAMBADA & MMLU Benchmark Evaluation Report

Evaluating small language models on long-range word prediction (LAMBADA) and multi-subject knowledge & reasoning (MMLU)

Submitted to: Prof. Anna Corazza
Submitted by: Francesco Ventimiglia, Danilo Rodriguez, Rohan Baidya
Github link: <https://github.com/ronvoy/gen-ai>
Site Link: <https://unina.cc/gen-ai>

1. Introduction

- We test three small language models on two complementary benchmarks.
- LAMBADA asks the model to predict the final word of a passage, guessable only from the full context - a focused test of long-range context handling.
- MMLU (Massive Multitask Language Understanding) is a multiple-choice knowledge and reasoning test spanning 57 subjects in four categories (STEM, Humanities, Social Sciences, Other); the model must write short step-by-step reasoning before committing to an answer letter.
- All models run through the OpenRouter API for a hardware-neutral comparison.
- Both benchmarks share one web app with live terminal output, run history, presets, and charts.

2. Small Language Models Used

2.1 Models at Glance

#	Model	Developer	Parameters	OpenRouter id	Architecture	Key Technique
1	Gemma-3-4B	Google	4B	google/gemma-3-4b-it	Dense decoder-only transformer with interleaved local/global attention	Knowledge distillation + local/global attention interleaving
2	Llama-3.2-3B	Meta	3B	meta-llama/llama-3.2-3b-instruct	Dense decoder-only transformer with Grouped Query Attention	Compact dense transformer
3	Ministral-8B	Mistral AI	8B	mistralai/ministral-8b-2512	Decoder-only transformer with Sliding	Sliding Window

#	Model	Developer	Parameters	OpenRouter id	Architecture	Key Technique
					Window Attention	Attention + GQA

2.2 Gemma-3-4B (Google)

- 4B instruction-tuned dense decoder-only transformer from Google, built with the same research that powers Gemini.
- Interleaves several local sliding-window attention layers with an occasional global attention layer, keeping memory low while information still flows across a long context (up to 128k tokens).
- The small Gemma 3 models are trained with knowledge distillation from larger teacher models.


Working technique - Knowledge distillation + local/global attention interleaving:

- Most layers only look at a local window of nearby tokens (cheap); every few layers one global layer lets any token attend to the whole context (expressive).
- Distillation teaches the small model to match a larger teacher's output distribution rather than raw text alone, delivering strong quality per parameter at edge-friendly cost.

Key properties:

- Dense decoder-only transformer, no expert routing.
- 5:1 interleaving of local sliding-window and global attention layers.
- Grouped Query Attention with QK-norm; 128k-token context window.
- Distilled from larger Gemma/Gemini-family teacher models.

Process flow - Gemma-3-4B inference



2.3 Llama-3.2-3B (Meta)

- 3B instruction-tuned model from Meta for on-device, low-cost use.
- Standard Llama recipe: dense transformer with RoPE, GQA, and SwiGLU layers.
- Built by pruning and distilling from larger Llama 3.1 models.


Working technique - Compact dense transformer:

- Takes the proven dense transformer design and shrinks it.
- Predictable and easy to serve on modest hardware.

Key properties:

- Dense decoder-only transformer.
- RoPE position embeddings for length generalisation.
- Grouped Query Attention and SwiGLU layers.

Process flow - Llama-3.2-3B inference



2.4 Ministral-8B (Mistral AI)

- 8B model from Mistral AI, built for edge use.
- Interleaved sliding window attention keeps memory and compute low on long inputs.
- Grouped query attention further shrinks the KV cache.


Working technique - Sliding Window Attention + GQA:

- Each layer attends only to a local window instead of every token.
- Stacked layers carry context further than any single window.
- GQA shares key-value heads for fast decoding on long passages.

Key properties:

- Sliding window attention for local context.
- Grouped Query Attention for a smaller KV cache.
- Efficient long-context inference at the edge.

Process flow - Ministral-8B inference



3. LAMBADA Benchmark

3.1 Dataset Overview

- LAMBADA is drawn from the BookCorpus.

- Each passage is chosen so the final word is predictable from the full passage but not from the last sentence alone.

Splits

Split	File	Passages	Purpose
Test	lambada_test_plain_text.txt	5,153	Primary evaluation
Development	lambada_development_plain_text.txt	4,869	Validation and tuning
Control Test	lambada_control_test_data_plain_text.txt	5,000	Baseline, unfiltered
Rejected	rejected_plain_text.txt	11,941	Passages cut during curation
Training Novels	train-novels/ (16 genres)	2,662 novels	Pre-training material
Vocabulary	lambada-vocab-2.txt	112,746 entries	Reference vocabulary


Properties

Property	Value
Source corpus	BookCorpus (unpublished novels)
Language	English
Task type	Word prediction
Curation criterion	Target guessable from full context only
First published	ACL 2016 (Paperno et al.)

3.2 Metrics

Metric	Definition	Better
Exact-match accuracy	correct / total after lowercasing and stripping punctuation	Higher
Average response time	Mean wall-clock time per API call	Lower
Error rate	errors / total (timeouts, failures)	Lower
Throughput	total / total wall-clock time	Higher

3.3 Scoring Flow



3.4 Fine-Tuning Parameters

- The Fine Tune panel exposes decoding parameters; it does not change model weights.

- For single-word prediction, greedy decoding with more examples works best.

Parameter	Range	Effect
Temperature	0.0 - 2.0	Higher adds randomness; 0.0 is greedy and deterministic
Top-p	0.0 - 1.0	Nucleus sampling cutoff; lower keeps only the likeliest tokens
Max tokens	1 - 128	Cap on answer length; a single word needs very few
Few-shot examples	0 - 5	Worked examples added to the prompt to set the format
Frequency penalty	-2.0 - 2.0	Discourages repeating tokens
Presence penalty	-2.0 - 2.0	Encourages introducing new tokens

Presets (see Section 6): Optimal (greedy decoding, five worked examples - most reliable accuracy), Normal (balanced defaults), Best Performance (trimmed token and example budget for the fastest, cheapest runs).

3.5 Results

Latest saved metrics per model on the test split.


Model	Accuracy (%)	Correct	Total	Avg Latency (s)	Errors
Gemma-3-4B	22.0	11	50	0.579	0
Llama-3.2-3B	20.8	208	1000	0.577	0
Minstral-8B	38.1	381	1000	0.731	0

Best accuracy: Minstral-8B at 38.1 percent. Fastest: Llama-3.2-3B at 0.577 s per query, with Gemma-3-4B essentially tied at 0.579 s.

Charts


Accuracy per model:

Accuracy



Average response time per model:

Average response time (s)



3.6 SLM Web App

The LAMBADA run panel: model and sample-count inputs, Run Benchmark and Fine Tune (presets) buttons, the live terminal streaming per-sample processing output (green = correct prediction, red = wrong), and the metrics table below with the latest

saved results for all three models.

Small Language Model Benchmark

Pick a model and run the LAMBADA word-prediction task through the OpenRouter API. Results are saved to the results folder and shown below.

Model

Samples (1-1000)

Gemma-3-4B

50

Run Benchmark

Fine Tune

History

lambada — benchmark output

[22:25:36] [42/50] ✓ target="stripes" predicted="stripes" 0.468s | accuracy 21.4%

[22:25:37] [43/50] ✗ target="izzy" predicted="too" 0.509s | accuracy 20.9%

[22:25:37] [44/50] ✗ target="macklin" predicted="more" 0.431s | accuracy 20.5%

[22:25:38] [45/50] ✗ target="renna" predicted="you" 0.78s | accuracy 20.0%

[22:25:38] [46/50] ✗ target="driveway" predicted="path" 0.464s | accuracy 19.6%

[22:25:39] [47/50] ✗ target="chore" predicted="always" 0.431s | accuracy 19.1%

[22:25:39] [48/50] ✓ target="treat" predicted="treat" 0.431s | accuracy 20.8%

[22:25:40] [49/50] ✓ target="julia" predicted="julia" 0.432s | accuracy 22.4%

[22:25:40] [50/50] ✗ target="taxi" predicted="driver" 0.606s | accuracy 22.0%

[22:25:40] Done: Gemma-3-4B - accuracy 22.0% (11/50), avg 0.579s, 0 error(s).

Done. Gemma-3-4B: 22% accuracy over 50 samples, 0.579s average.


Metrics

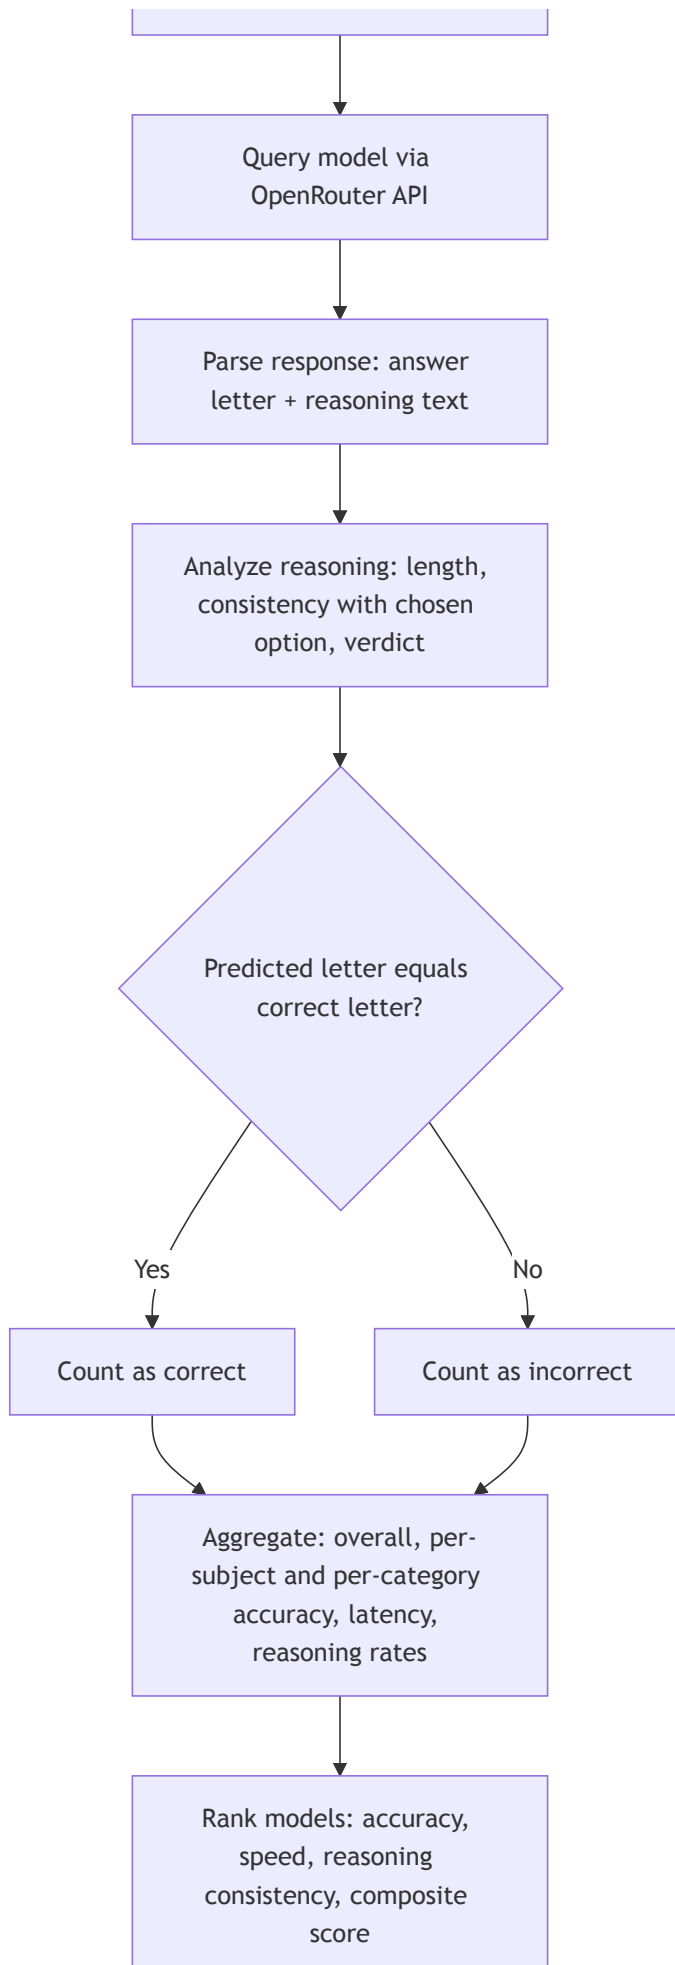
Model	Developer	Params	Accuracy (%)	Correct	Total	Avg time (s)	Errors
Gemma-3-4B	Google	4B	22	11	50	0.579	0
Llama-3.2-3B	Meta	3B	20.8	208	1000	0.577	0
Ministral-8B	Mistral AI	8B	38.1	381	1000	0.731	0

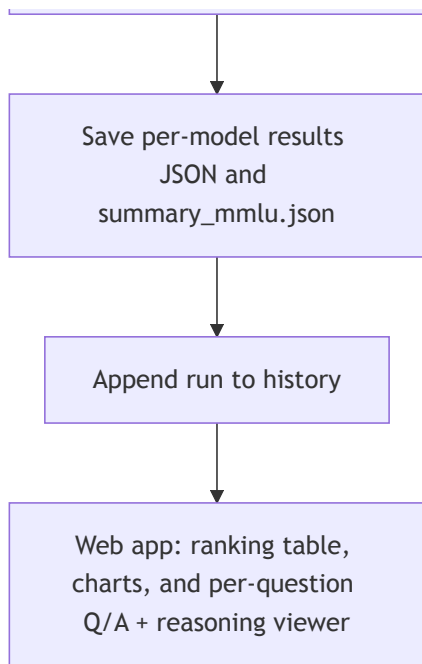
4. MMLU Benchmark

How the MMLU benchmark runs in this project, from a click on "Run MMLU" (or `./run_mmlu.sh`) to the ranking table, charts, and per-question reasoning shown in the web app. For every question the model must first write short step-by-step reasoning and then commit to one of four options (A-D); both the final letter and the reasoning text are parsed and evaluated.

4.1 Workflow Diagram







4.2 Chronological Steps

#	Step	What happens	Where (file / function)
1	Parameters	Read subjects, questions per subject, models, decoding params (all documented at the top of the script and shell runner)	evaluate_slm_mmlu.py (RUN PARAMETERS), run_mmlu.sh
2	Trigger	User clicks Run MMLU, runs <code>./run_mmlu.sh</code> , or <code>python evaluate_slm_mmlu.py</code>	passenger_wsgi.py (/mmlu/run) or run_mmlu_evaluation
3	Resolve subjects	Turn "all", a group preset, or a list into valid subject names	resolve_subjects
4	Fetch questions	Download the first 100 test rows per subject from the free Hugging Face datasets-server API (cais/mmlu, fallback tasksource/mmlu); no API key needed	_download_subject
5	Cache	Store rows in <code>_rsc/mmlu-dataset/<subject>.json</code> so reruns are offline and repeatable	fetch_subject_questions
6	Sample	Take the first N questions per subject (deterministic, no randomness)	fetch_subject_questions , load_mmlu_tasks
7	Build prompt	Chain-of-thought prompt: one worked example, strict Reasoning: then Answer: <letter> output format	build_mmlu_prompt , WORKED_EXAMPLE
8	Query model	Send the prompt to the chosen model through OpenRouter	query_model_mmlu (HTTP POST)
9	Parse	Extract the answer letter AND the reasoning text (handles Reasoning:/Answer: , <think> blocks, "the answer is (B)", bare letters, truncation)	parse_mmlu_response

#	Step	What happens	Where (file / function)
10	Evaluate reasoning	Score the reasoning: is it present, how long, does it actually support the chosen option; assign a verdict	analyze_reasoning
11	Compare	Exact match: predicted letter equals the correct letter	evaluate_model_mmlu
12	Aggregate	Overall / per-subject / per-category accuracy, latency, errors, reasoning rates	evaluate_model_mmlu
13	Rank	Per-dimension ranks and composite score across all evaluated models	build_mmlu_summary
14	Save	Write results/<model>_mmlu.json and results/summary_mmlu.json	run_mmlu_evaluation / run_mmlu_benchmark
15	History	Append the run (models, subjects, question count, params) to history	passenger_wsgi.append_history
16	Present	Ranking table, accuracy charts, and the per-question Q/A + reasoning accordion	templates/mmlu.html (Chart.js)

4.3 Components

Component	File	Role
Run parameters	evaluate_slm_mmlu.py (top), run_mmlu.sh (top)	Subject selection ("all", group presets, or explicit list), questions per subject (1-100), models, decoding params - every option listed in comments
Subject catalogue	MLLU_SUBJECTS , SUBJECT_GROUPS , CATEGORY_LABELS	The 57 subjects mapped to their official categories; group presets derived from the mapping
Dataset fetcher	fetch_subject_questions , _download_subject	Free HF datasets-server API client with source fallback and local JSON cache
Task loader	load_mmlu_tasks	Flattens subjects x questions into one ordered task list
Prompt builder	build_mmlu_prompt , WORKED_EXAMPLE	Chain-of-thought prompt with a worked example enforcing a parseable format
Model client	query_model_mmlu	OpenRouter chat-completions call with timing and error capture
Response parser	parse_mmlu_response	Splits a raw response into (answer letter, reasoning text)
Reasoning analyzer	analyze_reasoning	Heuristic quality check: presence, word count, consistency, verdict
Evaluator	evaluate_model_mmlu	Runs all tasks for one model; aggregates accuracy and reasoning metrics

Component	File	Role
Ranker	build_mmlu_summary	Per-dimension ranks + composite score across models
CLI runner	run_mmlu_evaluation , _parse_cli , run_mmlu.sh	One-shot terminal pipeline: setup, install, run, print ranking
Web routes	passenger_wsgi.py : /mmlu , /mmlu/run , /mmlu/metrics , /mmlu/details , /progress/<job_id>	Online runs, live progress feed, ranking JSON, and the per-question detail feed
Web page	templates/mmlu.html	Subject picker with presets, decoding sliders + presets, live terminal, ranking table, charts, Q/A + reasoning viewer

4.4 Metrics Produced

Metric	Definition
Accuracy	correct letters / total questions (exact match)
Category accuracy	Accuracy aggregated over STEM / Humanities / Social Sciences / Other
Subject accuracy	correct / total per individual subject
Average response time	Mean wall-clock seconds per API call
Errors	Count of failed or timed-out API calls
Reasoning rate	Share of answers that came with a non-trivial explanation (≥ 5 words)
Reasoning consistency	Share of answers whose reasoning actually supports the chosen option (names the letter or reuses the choice's key words)
Avg reasoning words	Mean length of the parsed reasoning text
Composite score	$0.70 \times \text{accuracy} + 0.15 \times \text{reasoning consistency} + 0.15 \times \text{relative speed}$ (fastest avg time / model's avg time)

4.5 Reasoning Verdicts (per question)

Verdict	Meaning
sound	Correct answer, and the reasoning clearly supports it
right-weak-link	Correct answer, reasoning present but does not clearly support the pick
lucky-guess	Correct answer with no real reasoning
flawed-reasoning	Reasoned its way to a wrong answer
blind-guess	Wrong answer and no reasoning

Verdict	Meaning
no-answer	No A-D letter could be parsed (API error or malformed output)

4.6 Results

Latest full run: all 57 subjects x 5 questions = 285 questions per model, zero API errors.

Rank	Model	Accuracy (%)	STEM	Human.	Social	Other	Reasoning consist. (%)	Avg words	Avg time (s)	Composite
1	Ministral-8B	78.9	78.9	73.9	83.3	80.0	69.1	54.1	1.537	0.731
2	Llama-3.2-3B	55.8	43.3	63.1	63.3	58.6	59.7	61.8	0.768	0.630
3	Gemma-3-4B	63.9	61.1	61.5	73.3	61.4	64.6	59.3	2.502	0.590


- Ministral-8B leads every category and the composite score.
- Gemma-3-4B is second on accuracy and reasoning consistency, but its 2.5 s average latency drops it below the much faster Llama-3.2-3B on the composite ranking.

Models are ranked per dimension (1 = best) on accuracy, speed, and reasoning consistency, then ordered by the composite score. Accuracy dominates the composite (70%); reasoning quality and latency (15% each) break ties, which mirrors how small language models are picked in practice: quality first, then cost and latency.

Charts


Overall accuracy per model, driving the composite ranking:

Accuracy by model



Subject-category accuracy (STEM / Humanities / Social Sciences / Other) per model:

Accuracy by category



4.7 MMLU Web App

The MMLU run panel: model and questions-per-subject inputs, subject picker with category presets, and the live terminal that streams processing output below the input fields before and during a run.

MMLU Benchmark - Knowledge & Reasoning

Multiple-choice questions across the 57 MMLU subjects, fetched live from the free Hugging Face datasets-server API. Each model must reason step by step before answering; the reasoning is parsed, scored for consistency, and shown per question below.

Model

Gemma-3-4B

Questions / subject (1-50)

3

Run MMLU

Decoding

Subjects (18 selected)

All 57

STEM

Humanities

Social Sciences

Other

Clear

☒ STEM (18)

☒ abstract algebra

☒ astronomy

☒ college biology

☒ college chemistry

☒ college computer science

☒ college mathematics

☒ college physics

☐ Humanities (13)

☐ formal logic

☐ high school european history

☐ high school us history

☐ high school world history

☐ international law

☐ jurisprudence

☐ logical fallacies

☐ Social Sciences (12)

☐ econometrics

☐ high school geography

☐ high school government and politics

☐ high school macroeconomics

☐ high school microeconomics

☐ high school psychology

☐ Other (14)

☐ anatomy

☐ business ethics

☐ clinical knowledge

☐ college medicine

☐ global facts

☐ human aging

☐ management

mmlu — benchmark output

\$ Ready. Pick subjects above and press Run MMLU — live processing output streams here.

The Q/A section lists every question the selected model saw, grouped per subject with a correct-count badge and category tag:

Questions, answers & model reasoning

Every question the selected model saw, its pick vs the correct answer, and the parsed reasoning behind it.

Gemma-3-4B

Load Q/A

Gemma-3-4B - 57 subject(s), 285 question(s). Reasoning given on 100.0% of answers, consistent with the chosen option on 64.6%.

Abstract Algebra	2/5 correct	STEM	
Anatomy	5/5 correct	Other	
Astronomy	4/5 correct	STEM	
Business Ethics	4/5 correct	Other	
Clinical Knowledge	3/5 correct	Other	
College Biology	5/5 correct	STEM	
College Chemistry	3/5 correct	STEM	
College Computer Science	3/5 correct	STEM	
College Mathematics	1/5 correct	STEM	
College Medicine	2/5 correct	Other	
College Physics	3/5 correct	STEM	
Computer Security	3/5 correct	STEM	

Expanding a subject shows each question with the model's pick vs the correct answer, its parsed reasoning, the reasoning verdict, and per-question latency:

Questions, answers & model reasoning

Every question the selected model saw, its pick vs the correct answer, and the parsed reasoning behind it.

Gemma-3-4B ▾

Load Q/A

Gemma-3-4B - 57 subject(s), 285 question(s). Reasoning given on 100.0% of answers, consistent with the chosen option on 64.6%.

Abstract Algebra 2/5 correct STEM

Q1. Find the degree for the given field extension $\mathbb{Q}(\sqrt{2}, \sqrt{3}, \sqrt{18})$ over $\mathbb{Q}$.

correct

Correct, weakly supported

4.547s

A. 0

B. 4 correct answer model's pick

C. 2

D. 6

Model reasoning:

We are looking for the degree of the field extension $\mathbb{Q}(\sqrt{2}, \sqrt{3}, \sqrt{18})$ over $\mathbb{Q}$. First, note that $\sqrt{18} = 3\sqrt{2}$, so $\mathbb{Q}(\sqrt{2}, \sqrt{3}, \sqrt{18}) = \mathbb{Q}(\sqrt{2}, \sqrt{3})$. The extension $\mathbb{Q}(\sqrt{2})$ over $\mathbb{Q}$ has degree 2, and the extension $\mathbb{Q}(\sqrt{2}, \sqrt{3})$ over $\mathbb{Q}(\sqrt{2})$ has degree 2. Therefore, the degree of the entire extension is $2 + 2 = 4$.

Q2. Let $p = (1, 2, 5, 4)(2, 3)$ in S_5 . Find the index of $\langle p \rangle$ in S_5 .

correct

Correct, weakly supported

8.615s

A. 8

B. 2

C. 24 correct answer model's pick

D. 120

Model reasoning:

First, we need to find the cycle type of the permutation $p = (1, 2, 5, 4)(2, 3)$. Applying the permutations from right to left, we have $(1, 2, 5, 4)(2, 3) = (1, 2, 5, 4)(2, 3) = (1, 2, 3, 5, 4)$. This is a 5-cycle. The order of a cycle of length n is n . Therefore, the order of p is 5. The index of the cyclic subgroup generated by p , denoted by $\langle p \rangle$, in the symmetric group S_5 is given by the order of the element p divided by their greatest common divisor. Since the order of p is 5, the index of $\langle p \rangle$ in S_5 is 5. However, this is not one of the answer choices. Let's recompute the permutation: $(1, 2, 5, 4)(2, 3) = (1, 2, 5, 4)(2, 3) = (1, 2, 3, 5, 4)$. The cycle decomposition is $(1)(2, 3, 5, 4)$. The order of the permutation is 5. The number of elements in S_5 is $5! = 120$. The index of $\langle p \rangle$ in S_5 is $120/5 = 24$.

5. Fine-Tuning Presets

Both benchmark pages expose one-click presets over the decoding parameters. They control decoding, not model weights.


5.1 LAMBADA Presets (PRESETS in config.py)

Preset	Temperature	Top-p	Max tokens	Few-shot	Intent
Optimal	0.0	1.0	16	5	Greedy decoding with the most worked examples - most reliable accuracy
Normal	0.3	0.9	32	3	Balanced default
Best Performance	0.0	1.0	8	2	Trimmed token and example budget - fastest, cheapest runs

5.2 MMLU Presets (MMLU_PRESETS in config.py)

Preset	Temperature	Top-p	Max tokens	Intent
Optimal	0.0	1.0	512	Most reasoning room - most reliable accuracy
Normal	0.2	0.95	384	Matches the defaults
Best Performance	0.0	1.0	192	Caps the reasoning budget so runs finish faster

6. Project Workflow



7. Conclusion

- Ministral-8B was the most accurate on both benchmarks (38.1% LAMBADA, 78.9% MMLU); Llama-3.2-3B was consistently the fastest.
 - Gemma-3-4B is competitive on accuracy and reasoning quality but pays for it in latency.
 - LAMBADA rewards real use of context over surface patterns; MMLU adds breadth of knowledge and rewards reasoning that actually supports the answer.
 - The usual trade-off holds: stronger context handling and reasoning help accuracy, smaller models answer faster.
 - Live terminal streaming and the Optimal / Normal / Best Performance presets make runs observable in real time and easy to reproduce.
-
- Paperno, D. et al. (2016). The LAMBADA dataset. Proceedings of ACL 2016.
 - Hendrycks, D. et al. (2021). Measuring Massive Multitask Language Understanding. Proceedings of ICLR 2021.
 - Google (2025). Gemma 3 technical report.
 - Meta (2024). Llama 3.2 model card.
 - Mistral AI (2024). Ministral model family.